

Apache Spark in CSB120

Installation Guide

Computer Science Department
Colorado State University

Download latest Apache Spark 3.1.2 binary (pre-build for Hadoop 3.2 and later): Go through its official documentation: <https://spark.apache.org/docs/latest/>

STEP 0: Your Hadoop cluster should be installed and running. To install Hadoop, please visit the relevant guide and follow the instructions and video clips in infospaces.

STEP 1: Downloading Spark

Log into your CS account. Download Spark release 3.1.2, pre-built for Hadoop 3.2 and later using following command.

```
wget https://dlcdn.apache.org/spark/spark-3.1.2/spark-3.1.2-bin-hadoop3.2.tgz
```

Unzip downloaded archive to your home directory. Change directory (cd) to folder containing downloaded archive (~/.spark-3.1.2-bin-hadoop3.2). Run following command to unzip archive and copy to your home directory.

```
tar -xvf spark-3.1.2-bin-hadoop3.2.tgz -C ~
```

STEP 2: Environmental variables

Set SPARK_HOME environment variable to point above unzipped directory.

Open .bashrc file, which will be inside your home directory. Add following line, if it doesn't exist.

```
export $SPARK_HOME=${HOME}/spark-3.1.2-bin-hadoop3.2
```

Reflect changes made in .bashrc file

```
source ~/.bashrc
```

Check the SPARK_HOME environment variable if it is correctly specified.

```
echo $SPARK_HOME
```

Note: Placeholders are denoted using <> in this document. You need to replace those (including <>) with appropriate values.

STEP 3: Configuration

Changes to configure Spark need to be done in \$SPARK_HOME/conf folder, which is inside unzipped folder (spark-3.1.2-bin-hadoop3.2). Templates for every configuration files are already given in this folder. We just need to make changes in required templates and save them with correct names.

workers

- This file stores the machines list/worker nodes (each machine name on new line) your Spark environment use to execute code.
- Save workers.template as workers.
- Remove localhost from the list and copy all the worker machines used in Hadoop (\$HADOOP_CONF_DIR/workers) and paste in workers.
 - This will maximize the benefit of data locality as same nodes are working as datanodes(HDFS) and workernodes(Spark).
- Save the `workers` file.

spark-env.sh

- This file stores primary configuration options for Spark environment.
- Spark provides several options for deployment, but we are configuring for Spark Standalone Deploy Mode. For more information, refer <http://spark.apache.org/docs/latest>
- Save spark-env.sh.template as spark-env.sh.
- Update following options (search for standalone) with appropriate values in spark-env.sh and don't forget to export them. You can refer sample below this table.

Environment variable	Description
SPARK_MASTER_IP	Bind the master to a specific hostname or IP address.
SPARK_MASTER_PORT	Start the master on a different port (default 7077). You need to select a set of available ports from the non-privileged port range. To reduce possible port conflicts, your will be each assigned a port range.
SPARK_MASTER_WEBUI_PORT	Port for the master webUI (default: 8080). Follow port selection instructions mentioned in SPARK_MASTER_PORT.
SPARK_WORKER_INSTANCES	To set the number of worker processes per node (e.g. 2, 4)
SPARK_WORKER_CORES	Total number of cores to allow Spark applications to use per worker instance (e.g. 1, 2) (default: all available cores)

SPARK_WORKER_MEMORY	Total amount of memory to allow Spark applications to use per worker instance, e.g. 1000m , 2g (default: total memory minus 1GB)
---------------------	--

Sample environment variables are given below. Put appropriate values in placeholders.

```
export SPARK_MASTER_IP=<hostname in CS120>
export SPARK_MASTER_PORT=<port number>
export SPARK_MASTER_WEBUI_PORT=<port number>
export SPARK_WORKER_CORES=2
export SPARK_WORKER_MEMORY=2g
export SPARK_WORKER_INSTANCES=2
```

Save the spark_env.sh file.

Spark-defaults.conf

- This file is used to set default properties included when running spark-submit. This is useful for setting default environmental settings.
- Each line in this file consists of a key and a value separated by whitespace.
- Save spark-defaults.conf.template as spark-defaults.conf.
- Update following options with appropriate values in spark-defaults.conf. You can refer sample below.

Sample key-value pairs are given below. Put appropriate values in placeholders.

```
spark.master spark://<SPARK_MASTER_IP>:<SPARK_MASTER_PORT>
```

Save the spark-defaults.conf file.

STEP 4: Launching Spark Cluster

- To be safe, start Hadoop cluster before launching Spark cluster.
- To start the cluster;
 - Login to Spark Masternode, specified in SPARK_MASTER_IP.
 - Run following script to launch Master as well as workers.

```
$SPARK_HOME/sbin/start-all.sh
```

- Check master webUI using <SPARK_MASTER_IP>:<SPARK_MASTER_WEBUI_PORT>
- To stop the cluster;
 - Login to Spark Masternode, specified in SPARK_MASTER_IP.
 - Run following script to launch Master as well as workers.

```
$SPARK_HOME/sbin/stop-all.sh
```

- You can individually start and stop master as well as worker instances. Refer <http://spark.apache.org/docs/latest/spark-standalone.html#cluster-launch-scripts> for additional information.

STEP 5: Launching Spark Applications

- Spark applications can be launched using spark-submit script.
- Change directory to your project folder.
- Run following command with appropriate values.

```
$SPARK_HOME/bin/spark-submit -class <your Class> --deploy-mode cluster -  
supervise <yourJar> <any_arguments>
```

You can refer <http://spark.apache.org/docs/latest/submitting-applications.html> for more information.